

Robert Davidson
ST1112: Statistics

70% Exam
30% Continuous Assessment (3 parts)

Contents

1	Descriptive Statistics	3
1.1	Sampling the mean	3
2	Interential Statistics - Interval Estimation	4
2.1	Confidence Intervals for a mean	4
2.2	Confidence Intervals for a population mean	6
2.3	Bootstrap, proportions and counts	7
2.3.1	Log Transform of the Data	7
2.3.2	Bootstrap	8

1 Descriptive Statistics

1.1 Sampling the mean

In **probability** we consider the underlying process which has some randomness or uncertainty, and we try to figure out what happens

In **statistics** we consider the data that we have, and we try to figure out what the underlying process is. The basic aim is to infer the population from the sample.

Example Consider a jar of red and green jelly beans

A probabilist starts by knowing the proportion of red and green jelly beans in the jar, and then tries to figure out the probability of drawing a red jelly bean.

A statistician starts by drawing a sample of jelly beans from the jar, and then tries to figure out the proportion of red and green jelly beans in the jar.

Definition : Central Limit Theorem

Sample means follow a normal distribution, centered on the population mean, with a standard deviation equal to population standard deviation divided by the square root of the sample size.

$$\bar{X} \sim N\left(\mu, \frac{\sigma}{\sqrt{n}}\right)$$

Definition : Standard Error

The standard error is the variability in the sampling distribution.

The standard error describes the typical difference between the sample measurement and the population parameter.

$$SE = \frac{\sigma}{\sqrt{n}}$$

Definition : Estimate σ

Often the value of the population standard deviation is unknown, and hence the standard error of the mean is unknown.

We can estimate the value of the standard error using the sample standard deviation (s) as an unbiased estimator of the population standard deviation (σ).

$$\sigma_{\bar{X}} = \frac{s}{\sqrt{n}}$$

2 Inferential Statistics - Interval Estimation

2.1 Confidence Intervals for a mean

Definition Confidence Interval for $n > 30$

For a large sample size, $n > 30$, a Confidence Interval for the population mean is given by:

$$\bar{X} \pm Z_{\frac{\alpha}{2}} \frac{s}{\sqrt{n}}$$

A 95% Confidence Interval, $\alpha = 0.05$, meaning we accept a 5% risk that our interval doesn't contain the true population mean.

This 5% is split into 2.5% in each tail of the distribution: $Z_{\frac{\alpha}{2}} = Z_{0.025}$.

When using a normal table that shows "area to the left", we need to find the Z value that corresponds to $1 - 0.025 = 0.975$. Thus: $Z_{0.025} = 1.96$.

95% Confidence is most commonly used because increasing the confidence level increases the width of the interval, this may not be useful.


Example The student newspaper wants to know how many students are exercising per week on average

- Take a **sample** from this **population**
- Estimate the **population parameter** using the **sample statistic**

```
st1112_data %>%
  select(exercise) %>%
  summarise(n = n(),
            mean = mean(exercise, na.rm=TRUE),
            sd = sd(exercise, na.rm=TRUE))

##      n mean  sd
## 1 54 6.19 4.11

st1112_data %>%
  ggplot(aes(x = exercise)) +
  labs(x = "Weekly exercise (hours)") +
  geom_density(colour = "blue",
              fill="blue", alpha=0.5) + theme_bw()
```



But a new survey on another 54 students would lead to a different estimate, so which should we report back to the newspaper? If we sample data from the population, there is uncertainty in our estimate of the population mean. The standard error of the mean is a measure of this uncertainty. In our example, the standard error of the mean is:

$$SE = \frac{4.11}{\sqrt{54}} \approx 0.6$$

We use the Central Limit Theorem to provide a range of values that will capture 95% of the sample means.

$$\bar{X} \pm Z_{\frac{\alpha}{2}} \times \frac{s}{\sqrt{n}} = 6.23 \pm 1.96 \times 0.6 = 5, 7.4$$

Example : Confidence intervals in R

```
st1112_data %>% select(exercise) %>% t.test()
##
##  One Sample t-test
##
## data:  .
## t = 11, df = 53, p-value = 2e-15
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
##  5.06 7.31
## sample estimates:
## mean of x
##      6.19
```

2.2 Confidence Intervals for a population mean

When σ is known, a 95% Confidence Interval for the population mean is given by:

$$\bar{X} \pm 1.96 \cdot \frac{\sigma}{\sqrt{n}}$$

When σ is unknown, a 95% Confidence Interval for the population mean is given by:

$$\bar{X} \pm 1.96 \cdot \frac{s}{\sqrt{n}}$$


But, what if σ is unknown and $n < 30$?

Definition Confidence Interval for $n < 30$

For a small sample size, $n < 30$, we use the **t-distribution instead of the normal distribution**. The t-distribution has heavier tails than the normal distribution and accounts for the additional uncertainty when estimating σ with s in small samples. The confidence interval is given by:

$$\bar{X} \pm t_{\frac{\alpha}{2}} \frac{s}{\sqrt{n}}$$

Where $t_{\frac{\alpha}{2}}$ is the critical value from the t-distribution with $n - 1$ **degrees of freedom**. The degrees of freedom represent the *number of independent pieces of information used to estimate the standard deviation*. As the degrees of freedom increase, the t-distribution approaches the normal distribution - infinite degrees of freedom is the normal distribution.



Example : Find t-value from tables For a 95% Confidence Interval

We have:

$$a = 1 - 0.95 = 0.05 \Rightarrow \frac{a}{2} = 0.025$$

We also have:

$$n = 13 \Rightarrow df = 13 - 1$$

We want to find the $t_{12,0.025}$ value from the t-distribution table, which is 2.179.

Example : Finding the t-value using R

```
qt(1 - 0.025, df = 12)
## [1] 2.179
```

2.3 Bootstrap, proportions and counts

What do we do when our data is not normally distributed? We have two options:


- Transform the data to make it normally distributed (e.g. log transformation, square root transformation)
- Use a non-parametric method (Bootstrap or CI for the population median)

2.3.1 Log Transform of the Data

Example : NBA Salaries


```
sals <- read.csv("data/NBA_season1718_salary.csv")

sals %>%
  filter(Tm=="DAL", season17_18>50000) %>%
  ggplot(aes(x=season17_18)) +
  geom_histogram() + xlab("Salary_($/year)")
```



The data is right-skewed, so we can apply a log transformation to make it more normally distributed.

```
sals %>%
  filter(Tm=="DAL", season17_18>50000) %>%
  ggplot(aes(x=log(season17_18))) +
  geom_histogram() + xlab("log(Salary)")
```



We can now calculate the mean and standard deviation of the log-transformed data.

```
sals %>% filter(Tm=="DAL", season17_18>50000) %>%
  mutate(log_sal = log(season17_18)) %>%
  select(log_sal) %>% summarise(n=n(),
                                mean=mean(log_sal),
                                sd=sd(log_sal))

##      n mean   sd
## 1  18 14.5 1.57

qt([p = 1 - 0.025, df = 18-1])
## [1] 2.11
```

Produce an interval estimate for the populat mean

```
14.5 - 2.11 * (1.57/sqrt(18))
14.5 + 2.11 * (1.57/sqrt(18))
## [1] 13.7 ## [1] 15.3
```

Take exponentials of the resulting interval to get back to original scale

```
exp(14.5 - 2.11 * (1.57/sqrt(18)))
exp(14.5 + 2.11 * (1.57/sqrt(18)))
## [1] 908172 ## [1] 4238843
```

2.3.2 Bootstrap

We can quantify the uncertainty in our estimate of the population mean by using the Central Limit Theorem or simulation via the Bootstrap method, as follows:

1. Take a bootstrap sample - random sample taken with replacement from the original sample (same size as original sample)
2. Calculate the bootstrap statistic - such as mean, median, proportion, etc.
3. Repeat steps 1 and 2 many times
4. Calculate the bounds of the 95% Confidence interval as the middle 95% of the bootstrap distribution